


CHAPTER 3: WINSOCK PROGRAMMING

Compiled by:
Madan Kadariya
NCIT



INTRODUCTION TO WINSOCK PROGRAMMING

The Windows Sockets Application Programming Interface (WinSock API) is a library of functions that implements the socket interface.

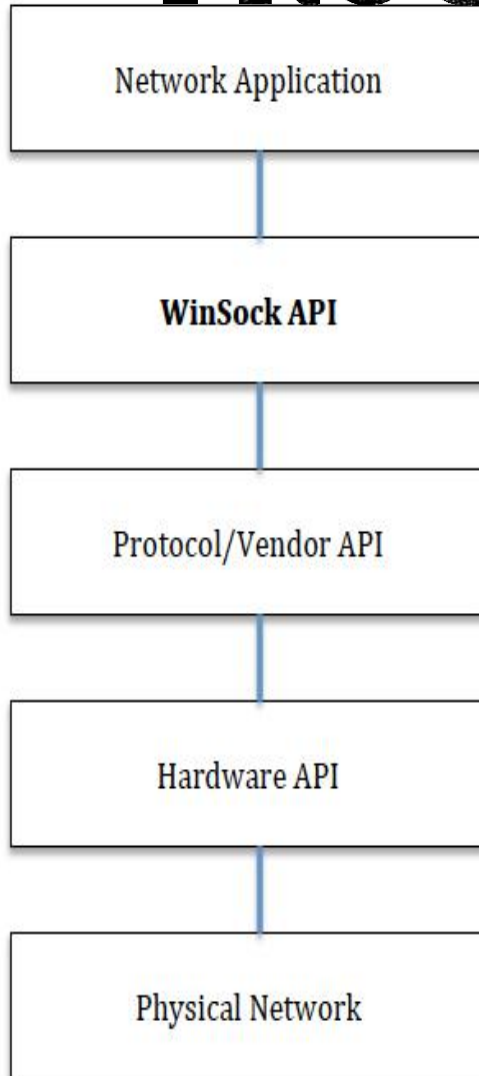
Winsock augments the Berkeley socket implementation by adding Windows-specific extension to support the message driven nature of the Windows operating system.

WinSock is a network application programming interface (API) for Microsoft Windows; a well-defined set of data structures and function calls implemented as a dynamic link library (DLL).

An application make function calls requesting generic network services (like send() and receive()); Winsock translates these generic requests into protocol- specific requests and performs the necessary tasks.

Residing between the application and the network implementation, Winsock shields you from the details of low-level network protocols.

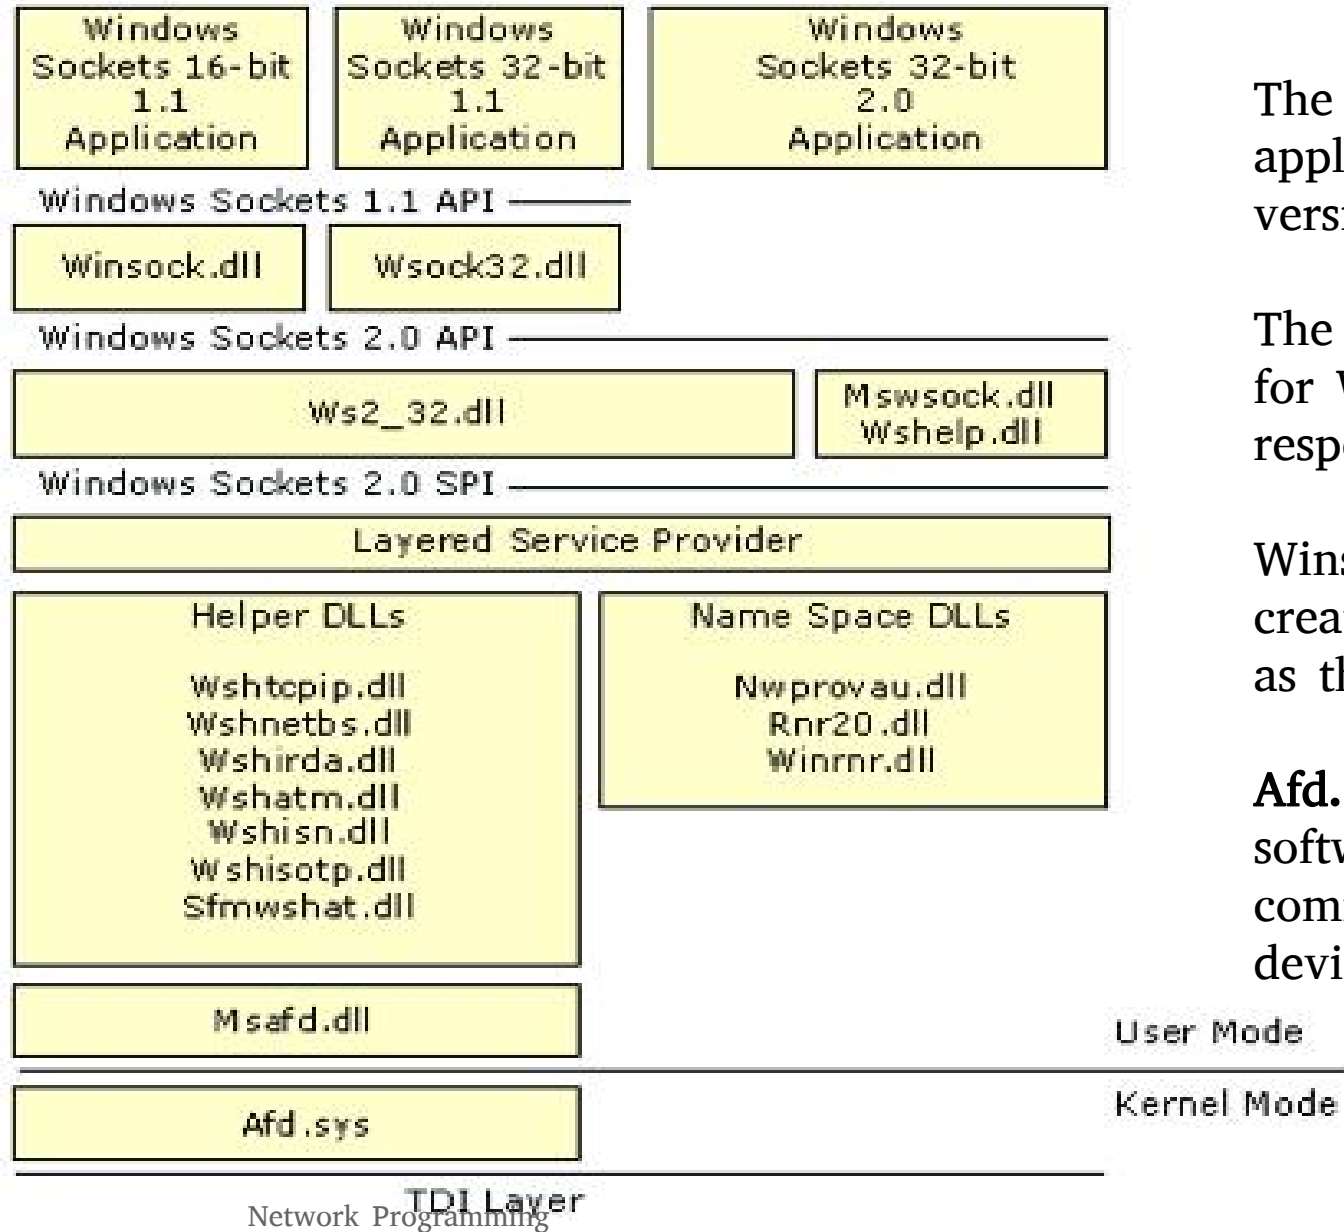
INTRODUCTION TO WINSOCK PROGRAMMING



The WinSock specification allows TCP/IP stack vendors to provide a consistent interface of their stacks so that application developers can write an application to the WinSock specification and have that application run on any vendor's WinSock-compatible TCP/IP protocol stack.

This is contrast to the days before the WinSock standard when software developers had to link their applications with libraries to each TCP/IP vendor's implementation. This limited the number of stacks that most applications ran on because of the difficulty in maintaining an application that used several different implementations of Berkeley sockets.

WINSOCK ARCHITECTURE



The top layer shows different network applications pertaining to different Winsock API versions and Windows architectures.

The two layers below it specify DLLs used for Winsock API version 1.1 and 2.0, respectively.

Winsock provides a Service Provider Interface for creating Winsock services, commonly referred to as the Winsock SPI.

Afd.sys is a Windows driver. A driver is a small software program that allows your computer to communicate with hardware or connected devices.

WINSOCK DLL (CURRENT IS WS2 32.DLL)

WINSOCK.DLL is a dynamic-link library that provides a common application programming interface (API) for developers of network applications that use the Transmission Control Protocol/Internet Protocol (TCP/IP) stack.

This means that a programmer who develops a Windows-based TCP/IP application, such as an FTP or Telenet client, can write one program that works with any TCP/IP protocol stack that provides Windows Socket Services (**WINSOCK.DLL**).

Some Winsock DLLs and their descriptions.

Winsock.dll

16-bit Winsock 1.1

Wsock32.dll

32-bit Winsock 1.1

Ws2 32.dll (this is latest)

Main Winsock 2.0

Mswsock.dll

Microsoft extensions to Winsock. Mswsock.dll is an API that supplies services that are not part of Winsock.

Ws2help.dll

Platform-specific utilities. Ws2help.dll supplies operating system-specific code that is not part of Winsock.

Wshtcpip.dll

Helper for TCP

Dynamic linking allows a module to include only the information needed to locate an exported DLL function at load time or run time. Dynamic linking differs from the more familiar static linking, in which the linker copies a library function's code into each module that calls it.

TYPES OF DYNAMIC LINKING

There are two methods for calling a function in a DLL:

In **load-time dynamic linking**, a module makes explicit calls to exported DLL functions as if they were local functions. This requires to link the module with the import library for the DLL that contains the functions. An import library supplies the system with the information needed to load the DLL and locate the exported DLL functions when the application is loaded.

In **run-time dynamic linking**, a module uses the **LoadLibrary** or **LoadLibraryEx** function to load the DLL at run time. After the DLL is loaded, the module calls the **GetProcAddress** function to get the addresses of the exported DLL functions. The module calls the exported DLL functions using the function pointers returned by **GetProcAddress**. This eliminates the need for an import library.

DLLS AND MEMORY MANAGEMENT

Every process that loads the DLL maps it into its virtual address space.

After the process loads the DLL into its virtual address, it can call the exported DLL functions.

The system maintains a per-process reference count for each DLL.

When a thread loads the DLL, the reference count is incremented by one. When the process terminates, or when the reference count becomes zero (run-time dynamic linking only), the DLL is unloaded from the virtual address space of the process.

Like any other function, an exported DLL function runs in the context of the thread that calls it. Therefore, the following conditions apply:

- The threads of the process that called the DLL can use handles opened by a DLL function. Similarly, handles opened by any thread of the calling process can be used in the DLL function.

- The DLL uses the stack of the calling thread and the virtual address space of the calling process.

- The DLL allocates memory from the virtual address space of the calling process.

ADVANTAGES OF DYNAMIC LINKING

Multiple processes that load the same DLL at the same base address share a single copy of the DLL in physical memory. Doing this saves system memory and reduces swapping.

When the functions in a DLL change, the applications that use them do not need to be recompiled or relinked as long as the function arguments, calling conventions, and return values do not change. In contrast, statically linked object code requires that the application be relinked when the functions change.

A DLL can provide after-market support. For example, a display driver DLL can be modified to support a display that was not available when the application was initially shipped.

Programs written in different programming languages can call the same DLL function as long as the programs follow the same calling convention that the function uses. The calling convention (such as C, Pascal, or standard call) controls the order in which the calling function must push the arguments onto the stack, whether the function or the calling function is responsible for cleaning up the stack, and whether any arguments are passed in registers.

A potential disadvantage to using DLLs is that the application is not self-contained; it depends on the existence of a separate DLL module. The system terminates processes using load-time dynamic linking if they require a DLL that is not found at process startup and gives an error message to the user. The system does not terminate a process using run-time dynamic linking in this situation, but functions exported by the missing DLL are not available to the program.

WINDOWS SOCKETS AND BLOCKING I/O

Berkeley Socket Versus WinSock

WinSock supports the TCP/IP domain for inter-process communication on the same computer as well as network communication. In addition to the TCP/IP domain, sockets in most UNIX implementations support the UNIX domain for the inter-process communication on the same computer.

The return values of certain Berkeley functions are different. For example, the **socket()** function returns **-1** on failure in the UNIX environment; the WinSock implementation returns **INVALID_SOCKET**.

Certain Berkeley functions have different names in WinSock. For example, in UNIX the **close()** system call is used to close the socket connection. In WinSock, the function is called **closesocket()**. It is so because WinSock socket handles may not be UNIX-style file descriptors.

WinSock Extension to Berkeley Sockets

WinSock has several extensions to Berkeley sockets. Most of these extensions are due to the message-driven architecture of Microsoft Windows. Some extensions are also required to support the non-preemptive nature of the 16-bit Windows operating environment.

Blocking I/O

The simplest form of I/O in Windows Sockets 2 is blocking I/O. Sockets are created in blocking mode by default.

Any I/O operation with a blocking socket will not return until the operation has been fully completed.

Thus, any thread can only execute one I/O operation at a time. For example, if a thread issues a receive operation and no data is currently available, the thread will block until data becomes available and is placed into the thread's buffer. Although this is simple, it is not necessarily the most efficient way to do I/O.

WinSock Asynchronous Functions

WinSock was originally designed for the non-preemptive Windows architecture. For this reason, several extensions were added to traditional Berkeley sockets.

Blocking Versus Non-blocking

Many of the Berkeley socket functions take an indeterminate amount of time to execute. When a function exhibits this behavior, it is said to block; calling the function blocks the further execution of the calling program.

In the Berkeley UNIX environment, for which sockets were originally developed, this didn't pose a serious problem because the UNIX operating system would simply preempt the blocking program and begin running another program.

Windows (unlike Windows NT) can't preempt a task, so all other programs are put on hold until the blocking call returns.

The designers of WinSock knew this posed a serious problem, so they added special code in the blocking functions to force the message loop of other applications to be checked. But this was still not the most efficient technique.

Berkeley sockets already have the notion of blocking versus non-blocking for some operations. For example, the `send()` function used to send data to a remote host may not return immediately, so the programmer is given the option of creating the socket with blocking or non-blocking sends.

If the socket is created in blocking mode, it won't return until the data has been delivered. If it's created in non-blocking mode, the call to the `send()` function returns immediately and the program must call another function called `select()` to determine the status of the send.

Windows and Windows NT can use the `select()` method of non-blocking calls, but the best thing to do in a Windows program is to use the special Windows asynchronous functions.

The special Windows asynchronous functions begin with the prefix **WSAAsync**. These functions were added to WinSock to make Berkeley sockets better fit the message-driven paradigm of Windows.

The most common events to use the asynchronous functions for are the sending and receiving of data. Sending data might not happen instantly, and receiving data most certainly will cause a program to wait unless it is receiving a constant stream of bytes.

By creating a socket for nonblocking sends and receives and using the **WSAAsyncSelect()** function call, an application will receive event notification messages to inform it when it can send data or when data has arrived and needs to be read. In the mean time, when there is no data communications occurring, the rest of the program remains fully responsive to the user's actions.

WinSock even extends Berkeley's nonblocking support to functions that could still cause a Berkeley UNIX program to block.

The name server's job is to take as input the plain text representation of a computer's name and return that computer's IP address.

GetXbyY is used to refer to database functions; function names take the form of get X by Y, or put another way: "Given Y, what is the corresponding X?"

gethostbyname(): given the computer's name, what is its host information? In Berkeley UNIX, the getXbyY functions may block.

WinSock adds asynchronous versions of the getXbyY functions called **WSAAsyncGetXbyY**. The gethostbyname() function is complimented by the nonblocking WinSock function called **WSAAsyncGetHostByName()**, for example. A call to a WSAAsyncGetXbyY function returns immediately with an identifying handle. When the actual work performed by the function has completed, a message is sent to the application notifying completion of the function with the specified handle.

The WinSock asynchronous functions were added primarily for the benefit of the nonpreemptive Windows environment.

The **WSAAsync** functions have an important use even in Windows NT. They allow an applications to remain responsive to the user.

Users won't enjoy working with any program if it forces them to wait for completion of a long event.

Most users expect a way to cancel operations that take a long time. For example, suppose that you have a program that takes as input a computer's plain-text name. The user enters the name and then presses a button labeled Look Up, which causes the `gethostbyname()` function to be called. Using `gethostbyname()` will cause the program to hang for an indeterminate amount of time until the request is carried out.

Under Windows NT other programs would still run, but under Windows the performance of all programs would be degraded. This program could be modified to use **WSAAsyncGetHostByName()** instead of `gethostbyname()`. As soon as users press the Look Up button, the **WSAAsyncGetHostByName()** function is called and returns an identifying handle.

If users wish to cancel the search, they can press the Cancel button, which terminates the request with that identifying handle. Users would maintain full control instead of being at the mercy of the program.

WINDOWS SOCKET EXTENSION; SETUP AND CLEANUP FUNCTION

The WinSock functions the application needs are located in the dynamic library named **WINSOCKET.DLL** or **WSOCK32.DLL** depending on whether the 16-bit or 32-bit version of Windows is being targeted.

The application is linked with either **WINSOCKET.LIB** or **WSOCK32.LIB** as appropriate.

The include file where the WinSock functions and structures are defined is named **WINSOCKET.H** for both the 16-bit and 32-bit environments.

Before the application uses any WinSock functions, the application must call an initialization routine called **WSAStartup()**.

Before the application terminates, it should call the **WSACleanup()** function.

WSAStartup

The **WSAStartup()** function initializes the underlying Windows Sockets Dynamic Link Library (WinSock DLL).

The **WSAStartup()** function gives the TCP/IP stack vendor a chance to do any application-specific initialization that may be necessary.

WSAStartup() is also used to confirm that the version of the WinSock DLL is compatible with the requirements of the application.

The prototype of the **WSAStartup()** functions follows:

```
int WSAStartup(WORD wVersionRequired, LPWSADATA lpWSADATA);
```

The **wVersionRequired** parameter is the highest version of the WinSock API the calling application can use. The high-order byte specifies the minor version and the low-order byte specifies the major version number. The **lpWSADATA** parameter is a pointer to a **WSADATA** structure that receives details of the WinSock implementation.

WinSock implementation.

```
typedef struct WSADATA {  
    WORD wVersion; // version supplied  
    WORD wHighVersion; // highest version that can be supplied  
    char  szDescription[WSADESCRIPTION_LEN+1]; /* a null terminated  
        ASCII string that describes the WinSock implementation*/  
    char  szSystemStatus[WSASYS_STATUS_LEN+1]; /* a null terminated  
        ASCII string that contains relevant status or configuration information */  
    unsigned short iMaxSockets; /* the maximum number of sockets that a  
        single process can potentially open */  
    unsigned short iMaxUdpDg; /* the size, in bytes, of the largest UDP  
        datagram that can be sent or received */  
    char FAR * lpVendorInfo; // a pointer to a vendor specific data structures  
} WSADATA;
```

WSACLEANUP

The **WSACleanup()** function is used to terminate an application's use of WinSock.

For every call to **WSAStartup()** there has to be a matching call to **WSACleanup()**.

WSACleanup() is usually called after the application's message loop has terminated.

In an MFC application, the **ExitInstance()** member function of the **CWinApp** class provides a convenient location to call **WSACleanup()**.

The prototype follows:

```
int PASCAL FAR WSACleanup(void);
```

WSA_GET_LAST_ERROR

The `WSA_Get_Last_Error()` function doesn't deal exclusively with startup or shutdown procedures, but it needs to be addressed early. Its function prototype looks like

```
int WSA_Get_Last_Error(void);
```

`WSA_Get_Last_Error()` returns the last WinSock error that occurred. Because WinSock isn't really part of the operating system but is instead a later addition, `errno` (like in UNIX and MS-DOS) couldn't be used.

As soon as a WinSock API call fails, the application should call `WSA_Get_Last_Error()` to retrieve specific details of the error.

FUNCTION FOR HANDLING BLOCKED I/O

The Berkeley method of using non-blocking sockets involves two functions: `ioctl()` and `select()`.

`ioctl()` is the UNIX function to perform input/output control on a file descriptor or socket.

Because a WinSock socket descriptor may not be a true operating system file descriptor, `ioctl()` can't be used, so `ioctlsocket()` is provided instead. `select()` is used to determine the status of one or more sockets. The use of `ioctlsocket()` to convert a socket to nonblocking mode looks like this:

```
// put socket s into nonblocking mode
```

```
u_long ulCmdArg = 1; // 1 for non-blocking, 0 for blocking  
ioctlsocket(s, FIONBIO, &ulCmdArg);
```

Once a socket is in its non-blocking mode, calling a normally blocking function simply returns `WSAEWOULDBLOCK` if the function can't immediately complete.

```
SOCKET clients;  
clients = accept(s, NULL, NULL);  
if (clients == INVALID_SOCKET)  
{  
    int nError = WSAGetLastError();  
    // if there is no client waiting to connect to this server,  
    nError will be WSAEWOULDBLOCK  
}
```

Your server application could simply call `accept()` periodically until the call succeeded, or you could use the `select()` call to query the status of the socket.

The `select()` function checks the readability, writeability, and exception status of one or more sockets.

WinSock provides a function called **WSAAsyncSelect()** to solve the problem of blocking socket function calls. It is a much more natural solution to the problem than using **ioctlsocket()** and `select()`.

It works by sending a Windows message to notify a window of a socket event. Its prototype is as follows:


```
int WSAAsyncSelect(SOCKET s, HWND hWnd, u_int wMsg, long lEvent);
```

s is the socket descriptor for which event notification is required. **hWnd** is the Window handle that should receive a message when an event occurs on the socket. **wMsg** is the message to be received by **hWnd** when a socket event occurs on socket **s**. It is usually a user-defined message (WM_USER + n). **lEvent** is a bitmask that specifies the events in which the application is interested.

WSAAsyncSelect() returns 0 (zero) on success and **SOCKET_ERROR** on failure. On failure, **WSAGetLastError()** should be called. **WSAAsyncSelect()** is capable of monitoring several socket events.

Event	Meaning
FD_READ	Socket ready for reading
FD_WRITE	Socket ready for writing
FD_OOB	Out-of-band data ready for reading on socket
FD_ACCEPT	Socket ready for accepting a new incoming connection
FD_CONNECT	Connection on socket completed
FD_CLOSE	Connection on socket has been closed

The **lEvent** parameter is constructed by doing a logical OR on the events in which you're interested. To cancel all event notifications, call **WSAAsyncSelect()** with **wMsg** and **lEvent** set to 0.

ASYNCHRONOUS DATABASE FUNCTION

WinSock provides a set of procedures commonly referred to as the database functions. The duty of these database functions is to convert the host and service names that are used by humans into a format usable by the computer.

The computers on an internetwork also require that certain data transmitted between them be in a common format. WinSock provides several conversion routines to fulfill this requirement.

Conversion Routines and Network Byte Ordering

There are four primary byte-order conversion routines. They handle the conversions to and from unsigned short integers and unsigned long integers.

Unsigned Short Integer Conversion

The `htons()` and `ntohs()` functions convert an unsigned short from host-to-network order and from network-to-host order. The prototype looks like.

```
u_short htons(u_short hostshort);  
u_short ntohs(u_short netshort);
```

Unsigned Long Integer Conversion

The `htonl()` and `ntohl()` functions work like `htons()` and `ntohs()` except that they operate on four-byte unsigned longs rather than unsigned shorts. The prototype look like the following:

```
u_long PASCAL htons(u_long hostlong);
```

```
u_long PASCAL ntohs(u_long netlong);
```

Converting IP Addresses

WinSock provides another set of conversion functions that provide a translation between the ASCII representation of a dotted-decimal IP address and the internal 32-bit, byte-ordered number required by other WinSock functions.

Converting an IP Address String to Binary

`inet_addr()` converts a dotted-decimal IP address string into a number suitable for use as an Internet address. Its function prototype is as follows:

```
unsigned long inet_addr(const char * cp);
```

`cp` is a pointer to a string representing an IP address in dotted-decimal notation. The `inet_addr()` function returns a binary representation of the Internet address given. This value is already in network byte order, so there is no need to call `htonl()`. If the `cp` string doesn't contain a valid IP address, `inet_addr()` returns `INADDR_NONE`. E.g.

```
u_long ulIPAddress = inet_addr("166.78.16.148");
```

Converting a Binary IP Address to a String

`inet_ntoa()` performs the opposite job of `inet_addr()`. Its function prototype is as follows:

```
char *  inet_ntoa(struct in_addr in);
```

in is a structure that contains an Internet host address. On success, the `inet_ntoa()` function returns a pointer to a string with a dotted-decimal representation of the IP address. On error, NULL is returned. A NULL value means that the IP address passed as the `in` parameter is invalid. E.g.

```
// first get an unsigned long with a valid IP address
```

```
u_long ulIPAddress = inet_addr("166.78.16.148");
```

```
// copy the four bytes of the IP address into an in_addr structure IN_ADDR in;
```

```
memcpy(&in, &ulIPAddress, 4);
```

```
// convert the IP address back into a string
```

```
char lpszIPAddress[16];
```

```
lstrcpy(lpszIPAddress, inet_ntoa(in));
```

What's My Name?

Some applications need to know the name of the computer on which they are running. The `gethostname()` function provides this functionality. The function's prototype looks like the following.

```
int gethostname(char * name, int namelen);
```

`name` is a pointer to a character array that will accept the null-terminated host name, and `namelen` is the size of that character array. The `gethostname()` function returns 0 (zero) on success and **SOCKET_ERROR** on failure. On a return value of **SOCKET_ERROR**, the application can call `WSAGetLastError()` to determine the specifics of the problem.

```
#define HOST_NAME_LEN 50
char lpszHostName[HOST_NAME_LEN]; // will accept the host name
char lpszMessage[100]; // informational message
if (gethostname(lpszHostName, HOST_NAME_LEN) == 0) {
printf(lpszMessage, "This computer's name is %s", lpszHostName);
} else {
    printf(lpszMessage, "gethostname() generated error %d", WSAGetLastError());
}
```

Finding a Host's IP Address

The function of **gethostbyname()** is to take a host name and return its IP address. This function, and its asynchronous counterpart named **WSAAsyncGetHostByName()**, may perform a simple table lookup on a host file local to the computer on which the program is running, or it may send the request across the network to a name server.

The function's prototype looks like the following:

```
struct hostent * gethostbyname(const char * name);
```

`name` is a pointer to a null-terminated character array that contains the name of the computer about which you want host information. The **hostent** structure returned has the following format:

```
struct hostent {  
    char * h_name; // official name of host  
    char ** h_aliases; // alias list  
    short h_addrtype; // host address type  
    short h_length; // length of address  
    char * * h_addr_list; // list of addresses  
    #define h_addr h_addr_list[0] // address, for backward compatibility  
};
```


Asynchronously Finding a Host's IP Address

In the `getXbyY` functions, one of which is `gethostbyname()`, the data retrieved might come from the local host or might come by way of a request over the network to a server of some kind.

Consequently, the application has to be concerned with response times and the responsiveness of the application to the user while those network requests are taking place.

The **`WSAAsyncGetHostByName()`** function is the asynchronous version of `gethostbyname()`.

The function prototype for `WSAAsyncGetHostByName()` is as follows:

```
HANDLE WSAAsyncGetHostByName(HWND hwnd, u_int wMsg, const char *name, char *buf, int buflen);
```

`hWnd` is the handle to the window to which a message will be sent when **`WSAAsyncGetHostByName()`** has completed its asynchronous operation.

`wMsg` is the message that will be posted to **`hWnd`** when the asynchronous operation is complete.

`name` is a pointer to a string that contains the host name for which information is being requested.

`buf` is a pointer to an area of memory that, on successful completion of the host name lookup, will contain the **`hostent`** structure for the desired host.

`buflen` is the size of the **`buf`** buffer. It should be `MAXGETHOSTSTRUCT` for safety's sake.

If the asynchronous operation is initiated successfully, the return value of **`WSAAsyncGetHostByName()`** is a handle to the asynchronous task. On failure of initialization, the function returns 0 (zero), and **`WSAGetLastError()`** should be called to find out the reason for the error.

Canceling an Outstanding Asynchronous Request

The handle returned by the asynchronous database functions, such as **WSAAsyncGetHostByName()**, can be used to terminate the database lookup.

The **WSACancelAsyncRequest()** function performs this task. Its prototype is the following:

```
int WSACancelAsyncRequest(HANDLE hAsyncTaskHandle);
```

hAsyncTaskHandle is the handle to the asynchronous task you wish to abort.

On success, this function returns 0 (zero).

On failure, it returns **SOCKET_ERROR**, and **WSAGetLastError()** can be called.

Finding a Host Name When You Know Its IP Address

The function of `gethostbyaddr()` is to take the IP address of a host and return its name. This function, and its asynchronous counterpart named `WSAAsyncGetHostByAddr()`, might perform a simple table lookup on a host file local to the computer on which the program is running, or it might send the request across the network to a name server. The function's prototype looks like the following:

```
struct hostent * PASCAL gethostbyaddr(const char * addr, int len, int type);
```

addr is a pointer to the IP address, in network byte order, of the computer about which you want host information.

len is the length of the address to which **addr** points. In WinSock 1.1, the length is always four because this version of the specification supports only Internet style addressing.

type must always be `PF_INET` for the same reason.

The `gethostbyaddr()` function returns a pointer to a `hostent` host entry structure on success and `NULL` on failure.

Upon a return value of `NULL`, you can call `WSAGetLastError()` to determine the specifics of the problem.

Asynchronously Finding a Host Name When You Know Its IP Address

The `WSAAsyncGetHostByAddr()` function is the asynchronous version of `gethostbyaddr()`. Its function prototype is as follows:

```
HANDLE WSAAsyncGetHostByAddr(HWND hWnd, u_int wMsg, const char * addr,
int len, int type, char * buf, int buflen);
```

hWnd is the handle to the window to which a message will be sent when `WSAAsyncGetHostByAddr()` has completed its asynchronous operation.

wMsg is the user defined message that will be posted to **hWnd** when the asynchronous operation is complete.

addr is a pointer to the IP address, in network byte order, of the computer about which you want host information.

len is the length of the address to which **addr** points and is always 4 (four) for Internet addresses. **type** must always be `PF_INET` because WinSock 1.1 supports only Internet-style addressing.

buf is a pointer to an area of memory that, upon successful completion of the address lookup, will contain the `hostent` structure for the desired host.

buflen is the size of the **buf** buffer. It should be `MAXGETHOSTSTRUCT` for safety's sake.

If the asynchronous operation is initiated successfully, the return value of `WSAAsyncGetHostByAddr()` is a handle to the asynchronous task. On failure of initialization, the function returns 0 (zero) and `WSAGetLastError()` should be called to find out the reason for the error.

Finding a Service's Port Number

The `getservbyname()` function gets service information corresponding to a specific service name and protocol. Its function prototype looks like the following:

```
struct servent * getservbyname(const char * name, const char * proto);
```

`name` is a pointer to a string that contains the service for which you are searching. `proto` is a pointer to a string that contains the transport protocol to use; it's either "udp", "tcp", or NULL. A NULL `proto` will match on the first service in the services table that has the specified name, regardless of the protocol. The `servent` structure returned has the following format:

```
struct servent {  
char * s_name; // official service name  
char ** s_aliases; // alias list  
short s_port; // port #  
char * s_proto; // protocol to use  
};
```

The `getservbyname()` function returns a pointer to a `servent` structure on success and NULL on failure. On a return value of NULL, you can call `WSAGetLastError()` to determine the specifics of the problem.

Asynchronously Finding a Service's Port Number

`WSAAsyncGetServByName()` is the asynchronous counterpart to `getservbyname()`. Its function prototype is as follows:

```
HANDLE WSAAsyncGetServByName(HWND hWnd, u_int wMsg, const char * name,  
const char * proto, char * buf, int buflen);
```

hWnd is the handle to the window to which a message will be sent when `WSAAsyncGetServByName()` has completed its asynchronous operation.

wMsg is the user defined message that will be posted to **hWnd** when the asynchronous operation is complete.

name is a pointer to a service name about which you want service information. **proto** is a pointer to a protocol name; it is “tcp”, “udp”, or NULL.

If **proto** is NULL, the first matching service is returned.

buf is a pointer to an area of memory that, on successful completion of the service lookup, will contain the servent structure for the desired service.

buflen is the size of the **buf** buffer.

It should be `MAXGETHOSTSTRUCT` for safety's sake.

Finding a Service Name When You Know Its Port Number

The `getservbyport()` function gets service information corresponding to a specific port and protocol. Its function prototype looks like the following:

```
struct servent FAR * PASCAL FAR getservbyport(int port, const char FAR *  
proto);
```

port is the service port, in network byte order.

proto is a pointer to a protocol name; it is “tcp”, “udp”, or NULL.

If **proto** is NULL, the first matching service is returned.

The `getservbyport()` function returns a pointer to a `servent` structure on success and NULL on failure.

On a return value of NULL, you can call `WSAGetLastError()` to determine the specifics of the problem.

Asynchronously Finding a Service Name When You Know Its Port Number

WSAAsyncGetServByPort() is the asynchronous counterpart to getservbyport(). Its function prototype is as follows:

```
HANDLE PASCAL FAR WSAAsyncGetServByPort(HWND hWnd, u_int wMsg, int port,
const char FAR * proto, char FAR * buf, int buflen);
```

hWnd is the handle to the window to which a message will be sent when WSAAsyncGetServByName() has completed its asynchronous operation.

wMsg is the user defined message that will be posted to hWnd when the asynchronous operation is complete.

port is the service port, in network byte order, of the service about which you want information. **proto** is a pointer to a protocol name; it is “tcp”, “udp”, or NULL.

If **proto** is NULL, the first matching service is returned.

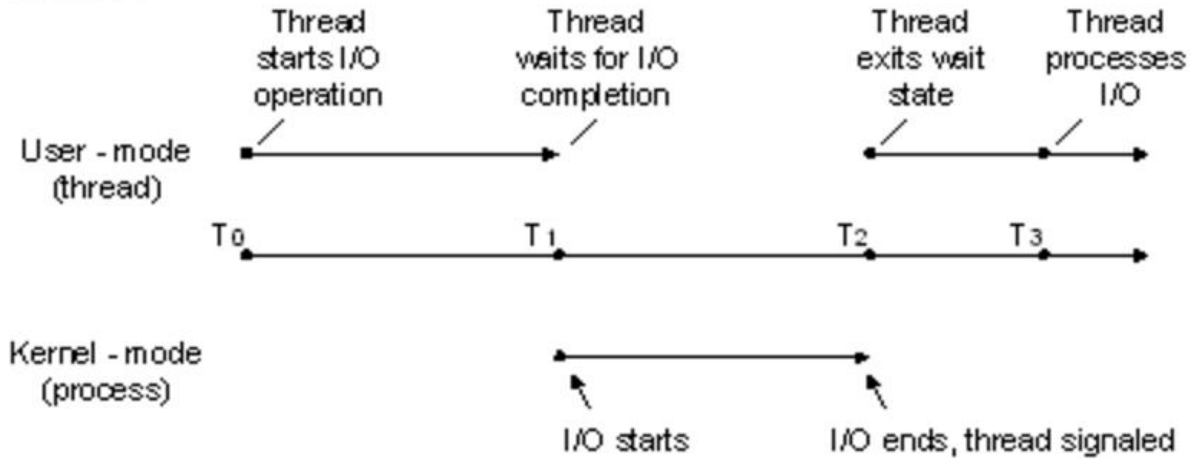
buf is a pointer to an area of memory that, on successful completion of the service lookup, will contain the servent structure for the desired service.

buflen is the size of the buf buffer. It should be MAXGETHOSTSTRUCT for safety's sake.

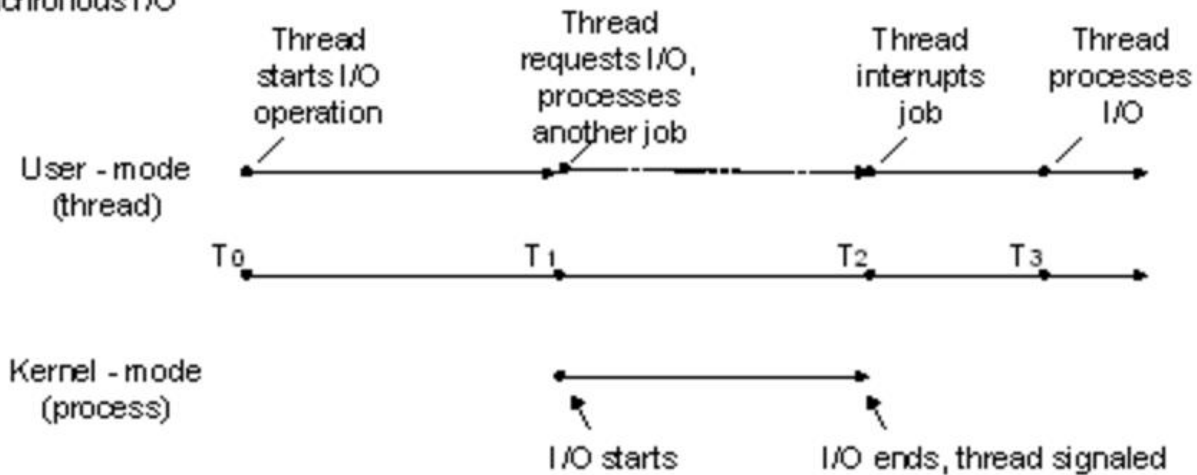
If the asynchronous operation is initiated successfully, the return value of WSAAsyncGetServByName() is a handle to the asynchronous task. On failure of initialization, the function returns 0 (zero) and WSAGetLastError() should be called to find out the reason for the error.

SYNCHRONOUS AND ASYNCHRONOUS I/O

Synchronous I/O



Asynchronous I/O



ASYNCHRONOUS I/O FUNCTIONS

There are two types of input/output (I/O) synchronization: **synchronous I/O** and **asynchronous I/O**.

Asynchronous I/O is also referred to as overlapped I/O.

In synchronous file I/O, a thread starts an I/O operation and immediately enters a wait state until the I/O request has completed.

A thread performing asynchronous file I/O sends an I/O request to the kernel by calling an appropriate function.

If the request is accepted by the kernel, the calling thread continues processing another job until the kernel signals to the thread that the I/O operation is complete.

It then interrupts its current job and processes the data from the I/O operation as necessary.

In situations where an I/O request is expected to take a large amount of time, such as a refresh or backup of a large database or a slow communications link, asynchronous I/O is generally a good way to optimize processing efficiency.

However, for relatively fast I/O operations, the overhead of processing kernel I/O requests and kernel signals may make asynchronous I/O less beneficial, particularly if many fast I/O operations need to be made. In this case, synchronous I/O would be better.

WinSock provides functions (**WSASend**, **WSASendTo**, **WSARecv**, **WSARecvFrom**, **WSAIoctl**, etc.) that support both synchronous and asynchronous I/O.

WSASocket is used to create overlapped socket.

```
SOCKET WSASocket (  
_In_ int          af,  
_In_ int          type,  
_In_ int          protocol,  
_In_ LPWSAPROTOCOL_INFO lpProtocolInfo,  
_In_ GROUP        g,  
_In_ DWORD        dwFlags);    // Used to create overlapped socket
```

af[in] : address family; AF_INET, AF_INET6, AF_UNSPEC

type[in] : type of the new socket; SOCK_STREAM, SOCK_DGRAM, SOCK_RAW

protocol[in] : the protocol to be used. If a value of 0 is specified, the caller does not wish to specify a protocol and the service provider will choose the protocol to use; else IPPROTO_TCP, IPPROTO_UDP, IPPROTO_ICMP

lpProtocolInfo[in] : A pointer to a WSAPROTOCOL_INFO structure that defines the characteristics of the socket to be created.

g [in] : An existing socket group ID.

dwFlags[in] : A set of flags used to specify additional socket attributes.

WSASEND

```
int WsaSend(  
_In_ SOCKET s,  
_In_ LPWSABUF lpBuffers,  
_In_ DWORD dwBufferCount,  
_Out_ LPDWORD lpNumberOfBytesSent,  
_In_ DWORD dwFlags,  
_In_ LPWSAOVERLAPPED lpOverlapped,  
_In_ LPWSAOVERLAPPED_COMPLETION_ROUTINE lpCompletionRoutine );
```

s[in] : A descriptor that identifies a connected socket.

lpBuffers[in] : A pointer to an array of WSABUF structures. Each WSABUF structure contains a pointer to a buffer and the length, in bytes, of the buffer.

dwBufferCount[in] : The number of WSABUF structures in the lpBuffers array.

lpNumberOfBytesSent[out] : The pointer to the number, in bytes, sent by this call if the I/O operation completes immediately. NULL for overlapped socket.

dwFlags[in] : The flags used to modify the behavior of the WsaSend function call.

lpOverlapped[in] : A pointer to a WSAOVERLAPPED structure. This parameter is ignored for non-overlapped sockets.

lpCompletionRoutine[in] : A pointer to the completion routine (function) called when the send operation has been completed. This parameter is ignored for non-overlapped sockets.

WSASENDTO

```
int WSA SendTo(  
_In_ SOCKET s,  
_In_ LPWSABUF lpBuffers,  
_In_ DWORD dwBufferCount,  
_Out_ LPDWORD lpNumberOfBytesSent,  
_In_ DWORD dwFlags,  
_In_ const struct sockaddr *lpTo,  
_In_ int iToLen,  
_In_ LPWSAOVERLAPPED lpOverlapped,  
_In_ LPWSAOVERLAPPED_COMPLETION_ROUTINE lpCompletionRoutine );
```

s[in] : A descriptor identifying a socket.

lpBuffers[in] : A pointer to an array of WSABUF structures. Each WSABUF structure contains a pointer to a buffer and the length of the buffer, in bytes. This array must remain valid for the duration of the send operation.

dwBufferCount[in] : The number of WSABUF structures in the lpBuffers array.

lpOverlapped[in] : A pointer to a WSAOVERLAPPED structure (ignored for nonoverlapped sockets).

lpCompletionRoutine[in] : A pointer to the completion routine called when the send operation has been completed (ignored for nonoverlapped sockets).

lpNumberOfBytesSent[out] : The pointer to the number, in bytes, sent by this call if the I/O operation completes immediately. NULL for overlapped socket.

dwFlags[in] : The flags used to modify the behavior of the WSA SendTo call.

lpTo[in] : An optional pointer to the address of the target socket in the SOCKADDR structure.

iToLen [in] : The size, in bytes, of the address in the lpTo parameter.

WSARECV

```
int WSARecv(  
_In_     SOCKET          s,  
_Inout_  LPWSABUF        lpBuffers,  
_In_     DWORD           dwBufferCount,  
_Out_    LPDWORD         lpNumberOfBytesRecv,  
_Inout_  LPDWORD         lpFlags,  
_In_     LPWSAOVERLAPPED lpOverlapped,  
_In_     LPWSAOVERLAPPED_COMPLETION_ROUTINE lpCompletionRoutine );
```

s[in] : A descriptor identifying a connected socket.

lpBuffers[in, out] : A pointer to an array of WSABUF structures. Each WSABUF structure contains a pointer to a buffer and the length, in bytes, of the buffer.

dwBufferCount[in] : The number of WSABUF structures in the lpBuffers array.

lpNumberOfBytesRecv[out] : A pointer to the number, in bytes, of data received by this call if the receive operation completes immediately. (NULL for overlapped socket)

lpFlags [in, out] : A pointer to flags used to modify the behavior of the WSARecv function call.

lpOverlapped[in] : A pointer to a WSAOVERLAPPED structure (ignored for non-overlapped sockets).

lpCompletionRoutine[in] : A pointer to the completion routine called when the receive operation has been completed (ignored for non-overlapped sockets).

WSARECVFROM

```
int WSARecvFrom(  
    _In_     SOCKET          s,  
    _Inout_  LPWSABUF        lpBuffers,  
    _In_     DWORD           dwBufferCount,  
    _Out_     LPDWORD         lpNumberOfBytesRecv,  
    _Inout_  LPDWORD         lpFlags,  
    _Out_     struct sockaddr *lpFrom,  
    _Inout_  LPINT            lpFromlen,  
    _In_     LPWSAOVERLAPPED lpOverlapped,  
    _In_     LPWSAOVERLAPPED_COMPLETION_ROUTINE lpCompletionRoutine  
);
```

s [in] : A descriptor identifying a socket.

lpBuffers [in, out] : A pointer to an array of WSABUF structures. Each WSABUF structure contains a pointer to a buffer and the length of the buffer.

dwBufferCount [in] : The number of WSABUF structures in the lpBuffers array.

lpNumberOfBytesRecv[out] : A pointer to the number of bytes received by this call if the WSARecvFrom operation completes immediately. (NULL for overlapped socket)

lpFlags [in, out] : A pointer to flags used to modify the behavior of the WSARecvFrom function call.

lpFrom [out] : An optional pointer to a buffer that will hold the source address upon the completion of the overlapped operation.

lpFromlen [in, out] : A pointer to a WSAOVERLAPPED structure (ignored for non-overlapped sockets).

lpCompletionRoutine [in] : A pointer to the completion routine called when the WSARecvFrom operation has been completed (ignored for nonoverlapped sockets).

WSAIoctl

```
int WSAIoctl(  
_In_ SOCKET s,  
_In_ DWORD dwIoControlCode,  
_In_ LPVOID lpvInBuffer,  
_In_ DWORD cbInBuffer,  
_Out_ LPVOID lpvOutBuffer,  
_In_ DWORD cbOutBuffer,  
_Out_ LPDWORD lpcbBytesReturned,  
_In_ LPWSAOVERLAPPED lpOverlapped,  
_In_ LPWSAOVERLAPPED_COMPLETION_ROUTINE lpCompletionRoutine );
```

s [in] : A descriptor identifying a socket.

dwIoControlCode [in] : The control code of operation to perform

lpInBuffer [in] : A pointer to the input buffer

cbInBuffer [in] : The size, in bytes, of the input buffer

lpOutBuffer [out] : A pointer to the output buffer

cbOutBuffer [in] : The size, in bytes, of the output buffer

lpcbBytesReturned [out] : A pointer to actual number of bytes of output

lpOverlapped [in] : A pointer to a WSAOVERLAPPED structure (ignored for non-overlapped sockets)

lpCompletionRoutine [in] : A pointer to the completion routine called when the operation has been completed (ignored for non-overlapped sockets)

ERROR HANDLING FUNCTIONS; ASYNCHRONOUS OPERATION

If asynchronous function returns `SOCKET_ERROR`, and the specific error code retrieved by calling `WSAGetLastError()` is `WSA_IO_PENDING`, then it means the overlapped operation has been successfully initiated and the completion will be indicated at a later time.

Any other error code indicates that the overlapped operation was not successfully initiated and no completion indication will occur.

When the overlapped operation completes the amount of data transferred is indicated either through the *cbTransferred* parameter in the completion routine(if specified), or through the *lpcbTransfer* parameter in `WSAGetOverlappedResult`.

BOOL WINAPI WSAGetOverlappedResult(In SOCKET s, In LPWSAOVERLAPPED lpOverlapped, Out LPDWORD lpcbTransfer, In BOOL fWait, Out LPDWORD lpdwFlags);

s [in]: A descriptor identifying the socket.

lpOverlapped [in]: A pointer to a `WSAOVERLAPPED` structure that was specified when the overlapped operation was started. This parameter must not be a `NULL` pointer.

lpcbTransfer [out]: A pointer to a 32-bit variable that receives the number of bytes that were actually transferred by a send or receive operation, or by the `WSAIoctl` function.

fWait [in]: A flag that specifies whether the function should wait for the pending overlapped operation to complete.

lpdwFlags [out]: A pointer to a 32-bit variable that will receive one or more flags that supplement the completion status.

If **WSAGetOverlappedResult** succeeds, the return value is **TRUE**. This means that the overlapped operation has completed successfully and that the value pointed to by *lpcbTransfer* has been updated.

If **WSAGetOverlappedResult** returns **FALSE**, this means that either the overlapped operation has not completed, the overlapped operation completed but with errors, or the overlapped operation's completion status could not be determined due to errors in one or more parameters to **WSAGetOverlappedResult**.

On failure, the value pointed to by *lpcbTransfer* will not be updated. Use **WSAGetLastError** to determine the cause of the failure (either by the **WSAGetOverlappedResult** function or by the associated overlapped operation).

USING NON-BLOCKING SOCKET, NON-BLOCKING WITH CONNECT

When connect is called on a blocking socket, the function blocks i.e. the function doesn't return until the TCP's 3-way handshake is completed (actually, until SYN, ACK is received from remote end).

The **ioctlsocket** function can be used to set socket in non-blocking mode. With a non-blocking socket, the connect returns immediately without completion.

In this case, connect will return `SOCKET_ERROR`, and `WSAGetLastError` will return `WSAEWOULDBLOCK`.

In this case, there are three possible scenarios.

Use the select function to determine the completion of the connection request by checking to see if the socket is writeable. If connect fails, failure of the connect attempt is indicated in `exceptfds` (application must then call `getsockopt` `SO_ERROR` to determine the error value to describe why the failure occurred).

If the application is using `WSAAsyncSelect` to indicate interest in connect events, then the application will receive an `FD_CONNECT` notification indicating that the connect operation is complete (successfully or not).

If the application is using `WSAEventSelect` to indicate interest in connection events, then the associated event object will be signaled indicating that the connect operation is complete (successfully or not).

SELECT IN CONJUNCTION WITH ACCEPT, SELECT WITH RECV/RECVFROM AND SEND/SENDTO

Select with accept

The accept function can block the caller until a connection is present if no pending connections are present on the queue, and the socket is marked as blocking.

If the socket is marked as non-blocking and no pending connections are present on the queue, accept returns an error.

When using non-blocking socket, select function can be used to determine if the connecting is present by checking to see if the socket is readable. Then, accept returns successfully, immediately.

Select with recv/recvfrom

With blocking socket, recv/recvfrom blocks if no data is available. When using non-blocking socket, select function can be used to determine if the data is available by checking to see if the socket is readable. Then, recv/recvfrom returns successfully, immediately.

Select with send/sendto

With blocking socket, send/sendto blocks if data can't be written to the kernel. When using non-blocking socket, select function can be used to determine if the data can be written to the kernel by checking to see if the socket is writable. Then, send/sendto returns successfully, immediately.

Getting Started with WinSock

General model for creating a streaming TCP/IP Server and Client.

Client

1. Initialize Winsock.
2. Create a socket.
3. Connect to the server.
4. Send and receive data.
5. Disconnect.

Server

1. Initialize Winsock.
2. Create a socket.
3. Bind the socket.
4. Listen on the socket for a client.
5. Accept a connection from a client.
6. Receive and send data.
7. Disconnect.

Initializing Winsock

All processes (applications or DLLs) that call Winsock functions must initialize the use of the Windows Sockets DLL before making other Winsock functions calls. This also makes certain that Winsock is supported on the system.

To initialize Winsock

1. Create a WSADATA object called wsaData.

```
WSADATA wsaData;
```

2. Call WSStartup and return its value as an integer and check for errors.

```
int iResult;
```

```
// Initialize Winsock
```

```
iResult = WSStartup(MAKEWORD(2,2), &wsaData);
```

```
if (iResult != 0) {
```

```
printf("WSStartup failed: %d\n", iResult);
```

```
return 1;
```

```
}
```

The **WSStartup** function is called to initiate use of WS2_32.dll. The WSADATA structure contains information about the Windows Sockets implementation. The MAKEWORD(2,2) parameter of WSStartup makes a request for version 2.2 of Winsock on the system, and sets the passed version as the highest version of Windows Sockets support that the caller can use.

CREATING A SOCKET FOR THE CLIENT

After initialization, a SOCKET object must be instantiated for use by the client.

To create a socket

1. Declare an **addrinfo** object that contains a **sockaddr** structure and initialize these values. For this application, the Internet address family is unspecified so that either an IPv6 or IPv4 address can be returned. The application requests the socket type to be a stream socket for the TCP protocol.

```
struct addrinfo *result = NULL, *ptr = NULL, hints;
```

```
ZeroMemory( &hints, sizeof(hints) );
```

```
hints.ai_family = AF_UNSPEC;
```

```
hints.ai_socktype = SOCK_STREAM;
```

```
hints.ai_protocol = IPPROTO_TCP;
```

2. Call the **getaddrinfo** function requesting the IP address for the server name passed on the command line. The TCP port on the server that the client will connect to is defined by DEFAULT_PORT as 27015 in this sample. The **getaddrinfo** function returns its value as an integer that is checked for errors.

```
#define DEFAULT_PORT "27015"  
// Resolve the server address and port  
iResult = getaddrinfo(argv[1], DEFAULT_PORT, &hints, &result);  
if (iResult != 0) {  
printf("getaddrinfo failed: %d\n", iResult);  
WSACleanup();  
return 1;  
}
```

3. Create a SOCKET object called ConnectSocket.

```
SOCKET ConnectSocket = INVALID_SOCKET;
```

4. Call the socket function and return its value to the ConnectSocket variable. For this application, use the first IP address returned by the call to getaddrinfo that matched the address family, socket type, and protocol specified in the hints parameter. In this example, a TCP stream socket was specified with a socket type of SOCK_STREAM and a protocol of IPPROTO_TCP. The address family was left unspecified (AF_UNSPEC), so the returned IP address could be either an IPv6 or IPv4 address for the server. If the client application wants to connect using only IPv6 or IPv4, then the address family needs to be set to AF_INET6 for IPv6 or AF_INET for IPv4 in the hints parameter.

```
// Attempt to connect to the first address returned by
```

```
// the call to getaddrinfo
```

```
ptr=result;
```

```
// Create a SOCKET for connecting to server
```

```
ConnectSocket = socket(ptr->ai_family, ptr->ai_socktype, ptr->ai_protocol);
```

5. Check for errors to ensure that the socket is a valid socket.

```
if (ConnectSocket == INVALID_SOCKET) {  
    printf("Error at socket(): %ld\n", WSAGetLastError());  
    freeaddrinfo(result);  
    WSACleanup();  
    return 1;  
}
```

If the socket call fails, it returns `INVALID_SOCKET`. The `if` statement in the previous code is used to catch any errors that may have occurred while creating the socket. `WSAGetLastError` returns an error number associated with the last error that occurred.

Connecting to a Socket

For a client to communicate on a network, it must connect to a server.

To connect to a socket

Call the connect function, passing the created socket and the sockaddr structure as parameters. Check for general errors.

// Connect to server.

```
iResult = connect( ConnectSocket, ptr->ai_addr, (int)ptr->ai_addrlen);
```

```
if (iResult == SOCKET_ERROR) {
```

```
    closesocket(ConnectSocket);
```

```
    ConnectSocket = INVALID_SOCKET;
```

```
}
```

// Should really try the next address returned by getaddrinfo

// if the connect call failed

```
freeaddrinfo(result);
```

```
if (ConnectSocket == INVALID_SOCKET) {
```

```
    printf("Unable to connect to server!\n");
```

```
    WSACleanup();
```

```
return 1;
```

Network Programming

```
}
```


The `getaddrinfo` function is used to determine the values in the `sockaddr` structure.

In this example, the first IP address returned by the `getaddrinfo` function is used to specify the `sockaddr` structure passed to the `connect`.

If the `connect` call fails to the first IP address, then try the next `addrinfo` structure in the linked list returned from the `getaddrinfo` function.

The information specified in the `sockaddr` structure includes:

- the IP address of the server that the client will try to connect to.

- the port number on the server that the client will connect to. This port was specified as port 27015 when the client called the `getaddrinfo` function.

SENDING AND RECEIVING DATA ON THE CLIENT

Client

```
#define DEFAULT_BUFLen 512

int recvbuflen = DEFAULT_BUFLen;

char *sendbuf = "this is a test";

char recvbuf[DEFAULT_BUFLen];

int iResult;

// Send an initial buffer

iResult = send(ConnectSocket, sendbuf, (int)
strlen(sendbuf), 0);

if (iResult == SOCKET_ERROR) {
printf("send failed: %d\n", WSAGetLastError());
closesocket(ConnectSocket);

WSACleanup();

return 1;

}

printf("Bytes Sent: %ld\n", iResult);
```

```
// shutdown the connection for sending since no more data will be sent
// the client can still use the ConnectSocket for receiving data
iResult = shutdown(ConnectSocket, SD_SEND);
if (iResult == SOCKET_ERROR) {
printf("shutdown failed: %d\n", WSAGetLastError());
closesocket(ConnectSocket);
WSACleanup();
return 1;
}
// Receive data until the server closes the connection
do {
iResult = recv(ConnectSocket, recvbuf, recvbuflen, 0);
if (iResult > 0)
printf("Bytes received: %d\n", iResult);
else if (iResult == 0)
printf("Connection closed\n");
else
printf("recv failed: %d\n", WSAGetLastError());
} while (iResult > 0);
```

Disconnecting the Client

Once the client is completed sending and receiving data, the client disconnects from the server and shutdowns the socket.

To disconnect and shutdown a socket

1. When the client is done sending data to the server, the shutdown function can be called specifying SD_SEND to shutdown the sending side of the socket. This allows the server to release some of the resources for this socket. The client application can still receive data on the socket.

// shutdown the send half of the connection since no more data will be sent

iResult = shutdown(ConnectSocket, SD_SEND);

if (iResult == SOCKET_ERROR) {

printf("shutdown failed: %d\n", WSAGetLastError());

closesocket(ConnectSocket);

WSACleanup();

return 1;

}

2. When the client application is done receiving data, the closesocket function is called to close the socket. When the client application is completed using the Windows Sockets DLL, the WSACleanup function is called to release resources.

```
// cleanup  
closesocket(ConnectSocket);  
WSACleanup();  
return 0;
```

CREATING A SOCKET FOR THE SERVER

After initialization, a SOCKET object must be instantiated for use by the server.

To create a socket for the server

1. The getaddrinfo function is used to determine the values in the sockaddr structure:

- AF_INET is used to specify the IPv4 address family.

- SOCK_STREAM is used to specify a stream socket.

- IPPROTO_TCP is used to specify the TCP protocol.

- AI_PASSIVE flag indicates the caller intends to use the returned socket address structure in a call to the bind function.

When the AI_PASSIVE flag is set and nodename parameter to the getaddrinfo function is a NULL pointer, the IP address portion of the socket address structure is set to INADDR_ANY for IPv4 addresses or IN6ADDR_ANY_INIT for IPv6 addresses.

27015 is the port number associated with the server that the client will connect to.

The addrinfo structure is used by the getaddrinfo function.

```
#define DEFAULT_PORT "27015"  
struct addrinfo *result = NULL, *ptr = NULL, hints;  
ZeroMemory(&hints, sizeof (hints));  
hints.ai_family = AF_INET;  
hints.ai_socktype = SOCK_STREAM;  
hints.ai_protocol = IPPROTO_TCP;  
hints.ai_flags = AI_PASSIVE;  
// Resolve the local address and port to be used by the server  
iResult = getaddrinfo(NULL, DEFAULT_PORT, &hints, &result);  
if (iResult != 0) {  
printf("getaddrinfo failed: %d\n", iResult);  
WSACleanup();  
return 1;  
}
```

2. Create a SOCKET object called ListenSocket for the server to listen for client connections.

```
SOCKET ListenSocket = INVALID_SOCKET;
```

3. Call the socket function and return its value to the ListenSocket variable. For this server application, use the first IP address returned by the call to getaddrinfo that matched the address family, socket type, and protocol specified in the hints parameter.

If the server application wants to listen on IPv6, then the address family needs to be set to AF_INET6 in the hints parameter.

If a server wants to listen on both IPv6 and IPv4, two listen sockets must be created, one for IPv6 and one for IPv4.

These two sockets must be handled separately by the application. Windows Vista and later offer the ability to create a single IPv6 socket that is put in dual stack mode to listen on both IPv6 and IPv4. For more information on this feature.

// Create a SOCKET for the server to listen for client connections

ListenSocket = socket(result->ai_family, result->ai_socktype, result->ai_protocol);

4. Check for errors to ensure that the socket is a valid socket.

```
if (ListenSocket == INVALID_SOCKET) {  
    printf("Error at socket(): %ld\n", WSAGetLastError());  
    freeaddrinfo(result);  
    WSACleanup();  
    return 1;  
}
```


Binding a Socket

For a server to accept client connections, it must be bound to a network address within the system. The following code demonstrates how to bind a socket that has already been created to an IP address and port. Client applications use the IP address and port to connect to the host network.

To bind a socket

The `sockaddr` structure holds information regarding the address family, IP address, and port number. Call the `bind` function, passing the created socket and `sockaddr` structure returned from the `getaddrinfo` function as parameters. Check for general errors.

```
// Setup the TCP listening socket  
  
iResult = bind( ListenSocket, result->ai_addr, (int)result->ai_addrlen);  
  
if (iResult == SOCKET_ERROR) {  
  
printf("bind failed with error: %d\n", WSAGetLastError());  
  
freeaddrinfo(result);  
  
closesocket(ListenSocket);  
  
WSACleanup();  
  
return 1;  
  
}
```

Once the `bind` function is called, the address information returned by the `getaddrinfo` function is no longer needed. The `freeaddrinfo` function is called to free the memory allocated by the `getaddrinfo` function for this address information.

```
freeaddrinfo(result);
```

Listening on a Socket

After the socket is bound to an IP address and port on the system, the server must then listen on that IP address and port for incoming connection requests.

To listen on a socket

Call the listen function, passing as parameters the created socket and a value for the backlog, maximum length of the queue of pending connections to accept. In this example, the backlog parameter was set to SOMAXCONN. This value is a special constant that instructs the Winsock provider for this socket to allow a maximum reasonable number of pending connections in the queue.

```
if ( listen( ListenSocket, SOMAXCONN ) == SOCKET_ERROR ) {  
    printf( "Listen failed with error: %ld\n", WSAGetLastError() );  
    closesocket(ListenSocket);  
    WSACleanup();  
    return 1;  
}
```

Accepting a Connection

Once the socket is listening for a connection, the program must handle connection requests on that socket.

To accept a connection on a socket

1. Create a temporary SOCKET object called ClientSocket for accepting connections from clients.

SOCKET ClientSocket;

2. Normally a server application would be designed to listen for connections from multiple clients. For high-performance servers, multiple threads are commonly used to handle the multiple client connections. One programming technique is to create a continuous loop that checks for connection requests using the listen function. If a connection request occurs, the application calls the accept, AcceptEx, or WSAAccept function and passes the work to another thread to handle the request.

ClientSocket = INVALID_SOCKET;

// Accept a client socket

ClientSocket = accept(ListenSocket, NULL, NULL);

if (ClientSocket == INVALID_SOCKET) {

printf("accept failed: %d\n", WSAGetLastError());

closesocket(ListenSocket);

WSACleanup();

return 1;

}

3. When the client connection has been accepted, a server application would normally pass the accepted client socket to a worker thread or an I/O completion port and continue accepting additional connections.

There are a number of other programming techniques that can be used to listen for and accept multiple connections.

These include using the select or WSAPoll functions.

Note: On Unix systems, a common programming technique for servers was for an application to listen for connections. When a connection was accepted, the parent process would call the fork function to create a new child process to handle the client connection, inheriting the socket from the parent. This programming technique is not supported on Windows, since the fork function is not supported. This technique is also not usually suitable for high-performance servers, since the resources needed to create a new process are much greater than those needed for a thread.

RECEIVING AND SENDING DATA ON THE SERVER

To receive and send data on a socket

```
#define DEFAULT_BUFLen 512
char recvbuf[DEFAULT_BUFLen];
int iResult, iSendResult;
int recvbuflen = DEFAULT_BUFLen;
// Receive until the peer shuts down the
connection
do {
iResult = recv(ClientSocket, recvbuf, recvbuflen, 0);
if (iResult > 0) {
printf("Bytes received: %d\n", iResult);
// Echo the buffer back to the sender
iSendResult = send(ClientSocket, recvbuf, iResult,
0);
if (iSendResult == SOCKET_ERROR) {
printf("send failed: %d\n", WSAGetLastError());
closesocket(ClientSocket);
WSACleanup();
return 1;
}
```

```
printf("Bytes sent: %d\n", iSendResult);
} else if (iResult == 0)
printf("Connection closing...\n");
else {
printf("recv failed: %d\n", WSAGetLastError());
closesocket(ClientSocket);
WSACleanup();
return 1;
} while (iResult > 0);
```

The send and recv functions both return an integer value of the number of bytes sent or received, respectively, or an error.

DISCONNECTING THE SERVER

Once the server is completed receiving data from the client and sending data back to the client, the server disconnects from the client and shutdowns the socket.

To disconnect and shutdown a socket

1. When the server is done sending data to the client, the shutdown function can be called specifying SD_SEND to shutdown the sending side of the socket. This allows the client to release some of the resources for this socket. The server application can still receive data on the socket.

// shutdown the send half of the connection since no more data will be sent

iResult = shutdown(ClientSocket, SD_SEND);

if (iResult == SOCKET_ERROR) {

printf("shutdown failed: %d\n", WSAGetLastError());

closesocket(ClientSocket);

Compiled by Chandra Mohan Jayaswal

WSACleanup();

return 1;

}

2. When the client application is done receiving data, the closesocket function is called to close the socket. When the client application is completed using the Windows Sockets DLL, the WSACleanup function is called to release resources.

// cleanup

closesocket(ClientSocket);

WSACleanup();

return 0;

END OF CHAPTER 3